AI Foundations & Policy Compliance

AI Tools for Operations Leaders

Your AI Journey — 4-Week Overview

Week 1
AI Foundations & Policy

- Tesla AI Policy
- Bottle Rocket basics
- Prompt engineering intro
- Your first AI prompts

Week 2
Advanced Prompts & Data

- Multi-shot prompting
- Data analysis with AI
- Report generation
- Building prompt libraries

Week 3
Workflows & Automation

- Multi-step AI workflows
- Team prompt sharing
- Process optimization
- Agent design basics

Week 4
Mastery & Rollout

- Advanced agent creation
- Team AI rollout plan
- Measuring ROI
- Capstone project

16 total hours of training — from zero AI experience to confident AI-powered leader

AI Foundations & Policy Compliance

AI Tools for Operations Leaders

240 minutes | Area General Managers | Tesla Distribution Centers

After This Course, You Will Be Able To...

✓ Use Bottle Rocket confidently for any operations task involving Tesla Data

✓ Write effective prompts using the ROLE + CONTEXT + TASK + FORMAT formula

✓ Make instant, correct decisions about which AI tools are allowed at Tesla

✓ Verify AI outputs and catch hallucinations before they reach leadership

✓ Build a personal prompt library that saves 30+ minutes per day

✓ Train your associates on Tesla-approved AI tools and policy compliance

Today's Agenda

1 Why AI Matters for Operations (10 min)
2 Tesla AI Policy Deep Dive (20 min)
3 Policy Quiz & Scenarios (16 min)
4 Introduction to Bottle Rocket (15 min)
5 Prompt Engineering Basics (15 min)
6 Exercise 1: First Guided Prompt (20 min)
7 Exercise 2: Operations Prompts + Personas (25 min)
8 *BREAK (10 min)*
9 AI Models & Tesla Agents (15 min)
10 Exercise 3: Bottle Rocket Deep Workshop (30 min)
11 Quality Checking AI Outputs (15 min)
12 Exercise 4: Iterate, Verify & Improve (25 min)
13 Wrap-up, Homework & Next Steps (14 min)

116 min

Hands-on (48%)

124 min

Instruction (52%)

Total: 240 minutes (4 hours)

Build a Multi-Agent Solution for Your DC

Design, build, and present an AI-powered workflow that solves a real problem at your distribution center.

01 Multi-Agent System
Build a project using multiple AI agents that work together — Nabu agents, Bottle Rocket prompts, or custom workflows

02 Power Automate Integration
Connect your AI workflows with business systems using Power Automate — triggers, approvals, or scheduled runs

03 Business-Relevant Problem
Solve a real operations challenge — headcount, safety, throughput, vendor management, or associate training

04 ROI Justification
Quantify your impact with hard numbers — hours saved per week, cost reduction, error rate improvement, or efficiency gains

YOU DELIVER • Working multi-agent workflow • Power Automate integration • ROI report with real metrics • 5-min live demo to peers

Timeline: Introduced Week 1 → Built Weeks 2–3 → Presented Week 4

Scope: Individual or paired project

01

Why AI Matters
for Operations

10 minutes

72% The Shift Is Already Happening

Of organizations have already deployed AI in at least one business function — and **supply chain leads adoption**. The question isn't if your competitors are using AI. It's how far ahead they already are.

McKinsey Global Survey on AI, 2024

COMPANIES WAITING

2.5× more likely to fall behind on operational efficiency within 18 months
Deloitte Digital Transformation Study

67% of leaders who delayed AI adoption reported talent retention issues
LinkedIn Workforce Report, 2024

$4.2M average cost of delayed AI implementation per year in missed productivity
Accenture Technology Vision

COMPANIES LEADING

40% reduction in time spent on routine reporting and data analysis tasks
Harvard Business Review, 2024

3.2× higher employee satisfaction when AI handles tedious work
MIT Sloan / BCG Research

23% improvement in supply chain decision speed — catching issues earlier
Gartner Supply Chain Survey

Bottom line: Tesla's culture is built on Speed and Excellence. AI is how we **multiply both**.

60% Where Does Your Time Actually Go?

Of a manager's week is spent on tasks that **don't require human judgment** — data gathering, formatting reports, writing routine communications, and chasing information across systems.

Microsoft Work Trend Index, 2024

TIME DRAINS (REPETITIVE)

11 hrs per week spent on email — reading, writing, sorting, and searching for buried info
McKinsey Global Institute

8 hrs per week gathering data from multiple systems and formatting into reports
Forrester Research

5 hrs per week on status updates, meeting prep, and predictable documentation
Asana Anatomy of Work Index

HIGH-VALUE WORK (WHAT YOU'RE HIRED FOR)

Strategic Long-term planning, process improvements, cost optimization — the work that drives your 20% improvement goals

Leadership Coaching your team, handling exceptions, making judgment calls that only you can make

Innovation Identifying new opportunities, building relationships, solving unsolved problems

The math is simple: If AI saves you 7.5 hours per week, that's **390 hours per year** — nearly 10 full work weeks redirected to what actually matters.

5× AI in Supply Chain Operations

Real use cases from your daily work — not hypotheticals. These are tasks **you already do** that AI can accelerate, improve, or handle entirely.

Based on distribution center operations workflows

BEFORE AI (MANUAL APPROACH)

DAILY REPORTS
45 min
Pulling data from multiple systems, formatting slides, writing commentary on variances

EMAIL COMMUNICATION
2+ hrs
Drafting updates to stakeholders, responding to inquiries, following up on exceptions

SOP DOCUMENTATION
3-4 hrs
Writing new procedures from scratch, updating existing docs, formatting for consistency

DATA ANALYSIS
1+ hr
Identifying trends in DPTO, cost per line, labor productivity — manually reviewing spreadsheets

WITH AI (YOUR NEW WORKFLOW)

DAILY REPORTS
10 min
AI drafts the narrative, you review and approve. Focus on insights, not formatting.

EMAIL COMMUNICATION
30 min
AI generates first drafts for routine updates. You customize tone and add judgment calls.

SOP DOCUMENTATION
45 min
AI creates comprehensive first draft from your bullet points. You refine and validate.

DATA ANALYSIS
15 min
AI surfaces anomalies and patterns. You decide what actions to take.

This isn't about replacing you. It's about **amplifying your impact** — more time for strategic thinking, leadership, and decisions only you can make.

AI Is a Tool, Not a Replacement

AI handles high-volume, repetitive analytical work at scale

You provide strategic judgment, exception handling, and contextual decisions

Organizations see 20-30% efficiency gains with proper AI training

Employees using AI tools save an average of 7.5 hours per week

Our target: save 30+ minutes per day — redirected to strategic work

Common trap: expecting magic instead of treating AI as a skilled assistant

7.5 hrs

saved per week
with AI tools

30+ min

daily target
for this course

02

Tesla AI Policy Deep Dive

20 minutes

Three Tiers of AI Tools at Tesla

APPROVED
Tesla Data OK

Chat.BottleRocket (go.tesla.com/chat)

IT Assist (it.bottlerocket.tesla.com)

Employee Assist (hr.bottlerocket.tesla.com)

Energy Chats (energy.bottlerocket.tesla.com)

Nabu RAG Agents
(nabu.bottlerocket.tesla.com)

Nova (nova.tesla.com)

M365 Copilot Chat (corporate link only)

CONDITIONAL
NO Tesla Data

GitHub Copilot (code only, no Tesla data)

GitHub Copilot CLI

Power Automate AI (no OpenAI connector)

PROHIBITED
Never for Work

ChatGPT, Claude.ai, Gemini web

Perplexity, Grok via X/xAI

Apple Intelligence (corp devices)

AI meeting transcription

AI wearables (Ray-Ban Meta, etc.)

DeepSeek (China-hosted)

Approved AI Tools

Source: Tesla Confluence — Approved Software > AI Tools (InfoSec)

If not explicitly approved, it is NOT allowed.

ANNOTATIONS

1. **Key rule:** Not approved = not allowed
2. **Tesla Data** = anything not public
3. **Bottle Rocket** — Tesla Data OK
4. **IT Assist** — IT help desk bot
5. **Employee Assist** — HR bot

YOUR 3 SAFE TOOLS
Bottle Rocket
IT Assist
Employee Assist

Conditional & Restricted Tools

Same policy page — scroll down for the tricky rules

ANNOTATIONS

1. GitHub Copilot — approved but NO Tesla code
2. Copilot in Office — M365 is NOT available
3. M365 Chat — OK but use *corporate link only*
4. Teams Recording — needs policy approval, vendors blocked

COMMON MISTAKES

Using copilot.microsoft.com instead of m365.cloud.microsoft.chat

Letting vendors record or transcribe calls

When in doubt: Use Bottle Rocket. It is always approved for Tesla Data.

What Counts as "Tesla Data"?

Any information not publicly available outside Tesla

Employee names, headcount, schedules, shift assignments

KPIs, throughput rates, error rates, safety incident details

Vendor names, contract terms, customer order information

Internal SOPs, financial data, email content about Tesla business

X WRONG	✓ RIGHT
"Sort my DC's headcount data by shift" in ChatGPT	"Sort my DC's headcount data by shift" in Bottle Rocket

When in doubt, assume it is Tesla Data and use Bottle Rocket.

Tier 1: Sensitive Data

Highest risk — never share outside Tesla secure systems

BUSINESS DATA

Core Information Assets
· Trade secrets, · Source code, · Autopilot/CV code, · CAD/3D design files, · Code name projects, · M&A information

Non-Public Financial Data
· Financial statements not yet publicly reported

Security Business Data
· Service secrets, · Passwords, · Encryption keys, · Device fingerprinting

Energy Data
· Tesla IP & trade secrets, · Site controller & Megapack SW/FW, · RTAC proprietary program

PERSONAL DATA

Location & Government IDs
· GPS coordinates, · SSN, · Tax ID, · Passport, · Driver's license

Health & Biometric
· Employee health info, · Facial recognition, · Voice profiles, · Fingerprints

Financial & Transactions
· Credit card numbers, · Bank/routing numbers, · Transaction details

Behavioral & Demographic
· Customer behavior profiles, · Race, · Ethnicity, · Religion, · Sexual orientation, · Disability

Background Check Data
· Credit/background reports, · Criminal history

Source: Tesla Data Classification & Handling Policy

Tier 2: Confidential Business Data

Internal use only — approved tools required for AI processing

BUSINESS DATA

Core Info & Security
· Board meeting agendas, · non-sensitive source code, · JIRA tickets, · pen test reports, · risk assessments

Internal Corporate
· Org structure, · employee roles, · data center address, · third party contracts, · server logs

Operational Info
· Internal policies & procedures, · meeting notes, · staff presentations, · aggregated statistics, · Intranet content

Product & Manufacturing
· Design/manufacturing/testing info, · floor plans, · equipment specs, · design documents, · engineering docs

Energy & Aggregated
· SCADA configs, · network diagrams, · cyber asset lists, · aggregated location/device/transaction data

PERSONAL DATA

Customer Contact
· Name, · address, · phone, · email, · date of birth

Location & Transactions
· SDK/beacon/IP location data, · service records, · credit card details

Tesla Identifiers
· Employee ID, · vehicle pseudonym ID, · referral invite code

Customer Content & Comms
· Feedback/surveys, · forum content, · calls/texts/voice memos with customers

Device & Vehicle Data
· Device type, OS, carrier, app version; · license plate, VIN, proof of insurance

HR / Talent Data
· Contact details, · CV, · payroll, · employment history, · benefits, · performance data

Photos & Recordings
· Customer uploaded photos, · Autopilot camera recordings (safety events, FSD Beta)

Tier 3: Public Data

Publicly available — safe to use with any AI tool

BUSINESS DATA

- **Public Information**
  Published Tesla corporate info (job announcements), published financial information, public presentations

- **Press & Media**
  Press releases, blogs, social media posts

- **Regulatory Filings**
  Filings subject to public disclosure such as SEC filings

- **Marketing Materials**
  Published marketing materials or brochures

- **Product Specs**
  Megapack product specifications

- **Public Presentations**
  Conference talks, publicly shared decks

This is the only tier where external AI tools (ChatGPT, Claude, etc.) could theoretically be used — but Tesla policy **still prohibits them**. Use Bottle Rocket for everything.

The Grok Trap — Same Model, Different Rules

Grok via Bottle Rocket

Platform	go.tesla.com/chat
Tesla Data	YES — fully approved
Routing	Tesla secure infrastructure

Use freely with internal data

Grok via X / xAI Apps

Platform	x.com, xAI website, xAI apps
Tesla Data	NO — policy violation
Routing	External xAI servers

Never use for Tesla work

The platform determines data protection, not the model name.

The Two-Question Test

QUESTION 1
Is this tool listed at go.tesla.com/aitools?

QUESTION 2
Does my task involve Tesla Data?

YES + YES	YES + NO	NO to Q1
Use Approved Tool	**Approved or Conditional**	**STOP**
Bottle Rocket, Nova, IT/Employee Assist	Any listed tool is fine for non-Tesla data	Tool is prohibited for work use

Violations can result in disciplinary action up to and including termination.

INTERACTIVE QUIZ **Policy Rapid Fire**

#	TOOL / ACTION	ANSWER
1	Bottle Rocket	Approved
2	Grok via X app	Prohibited
3	Grok via Bottle Rocket	Approved
4	Claude.ai (website)	Prohibited
5	Claude inside Bottle Rocket	Approved
6	M365 Copilot (m365.cloud.microsoft.chat)	Approved
7	M365 Copilot (copilot.microsoft.com)	Prohibited
8	GitHub Copilot with Tesla code	Prohibited
9	GitHub Copilot with generic Python	Conditional
10	AI meeting transcription	Prohibited
11	Apple Intelligence on corp device	Prohibited
12	Perplexity for generic research	Prohibited
13	Nova searching Confluence	Approved
14	Employee Assist for HR	Approved
15	DeepSeek	Prohibited

5 MIN

rapid fire round

KEY TAKEAWAY

Same model can be **Approved** OR **Prohibited** depending on the platform.

Policy Rapid Fire

For each item, call out: APPROVED / CONDITIONAL / PROHIBITED

#	TOOL OR ACTION	ANSWER
1	Bottle Rocket	?
2	Grok via X app	?
3	Grok via Bottle Rocket	?
4	Claude ai (website)	?
5	Claude inside Bottle Rocket	?
6	M365 Copilot Chat (m365 cloud microsoft chat)	?
7	M365 Copilot Chat (copilot.microsoft.com)	?
8	GitHub Copilot with Tesla code	?
9	GitHub Copilot with generic Python code	?
10	AI meeting transcription (approved tool)	?
11	Apple Intelligence on personal iPhone at work	?
12	Perplexity AI for generic research	?
13	Nova searching Confluence	?
14	Employee Assist for HR questions	?
15	DeepSeek (China-hosted)	?

Green = Approved · Yellow = Conditional · Red = Prohibited

#	TOOL OR ACTION	STATUS
1	Bottle Rocket	APPROVED
2	Grok via X app	PROHIBITED
3	Grok via Bottle Rocket	APPROVED
4	Claude.ai (website)	PROHIBITED
5	Claude inside Bottle Rocket	APPROVED
6	M365 Copilot Chat (m365.cloud.microsoft.chat)	APPROVED
7	M365 Copilot Chat (copilot.microsoft.com)	PROHIBITED
8	GitHub Copilot with Tesla code	PROHIBITED
9	GitHub Copilot with generic Python code	CONDITIONAL
10	AI meeting transcription	PROHIBITED
11	Apple Intelligence on personal iPhone at work	OK
12	Perplexity AI for generic research	PROHIBITED
13	Nova searching Confluence	APPROVED
14	Employee Assist for HR questions	APPROVED
15	DeepSeek (China-hosted)	PROHIBITED

What Would You Do? — Manager Scenarios

The VP Request

Your VP messages at 4:45 PM: "I need a Dive Deep on Q4 throughput by 8 AM." A colleague says "use ChatGPT — it's faster."

Urgency doesn't override policy.

The Data Leak

Your supervisor uploaded weekly KPI spreadsheets to Perplexity for a month. They didn't know it was wrong.

Stop. Report. Rebuild on approved tools.

The Competitive Pressure

A peer AGM's team uses unapproved tools and gets better numbers. Do you match them?

Policy violations aren't justified by productivity.

7 MIN
table discussion

Discuss with your table — 7 minutes

Gray Zones — It's Trickier Than You Think

"I anonymized the data first"
Removing names doesn't make Tesla Data non-Tesla. Overtime patterns and staffing ratios from inside Tesla are still Tesla Data.

"It's my personal phone, on my own time"
The policy follows the DATA, not the clock or the device.

"I only shared the AI output"
AI output containing Tesla Data must be handled like any Tesla Data.

"I used Bottle Rocket first, then polished in Claude.ai"
Two-step workarounds add risk for zero benefit. Bottle Rocket runs Claude models — just refine there.

ⓘ When in doubt: use Bottle Rocket, keep it in approved channels, and ask.

03

Introduction to Bottle Rocket

15 minutes

Inside Bottle Rocket

Access at go.tesla.com/chat — sign in with Tesla credentials

Chat interface: type your prompt, AI responds in real-time

Model selector (top): choose from 8 AI models for different tasks

Personas (sidebar): pre-built AI personalities for specialized tasks

File upload: drag & drop PDFs, spreadsheets, images for AI analysis

Web search toggle (globe icon): ground responses with live internet data

Choosing Your AI Model

DEFAULT **Qwen 3 30B** Fast Daily tasks, quick answers	**Qwen 3 235B** Medium Complex analysis, reasoning	**OpenAI GPT OSS 120B** Medium Writing, general tasks	**Grok 4** Medium Deep research, nuanced analysis
Grok 4 Mini Fast Quick reasoning, fast iteration	**Grok 4 Mini (High)** Med-Fast Balanced Grok quality/speed	**Gemini 3 Pro** Medium Multimodal, image analysis	**Gemini 2.5 Pro** Slower Long docs, complex tasks

Start with the default (Qwen 3 30B). Switch models for deeper analysis.

Personas — Customize Your AI Assistant

Personas ∧

Tyler
Tyler is a dive deeper for t…

Feynman
This persona breaks down any …

General
Your reliable and knowledgeab…

The Swiss Army Knife
An AI assistant that acts like a
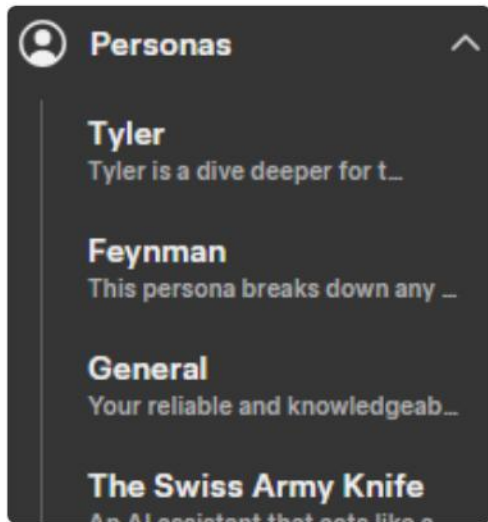
Bottle Rocket — Personas sidebar

🔴 **Feynman**
Explains complex topics in simple terms. Great for training materials and breaking down technical concepts.

⚙ **General**
Balanced all-purpose assistant. Good default for everyday operations tasks.

🔧 **Swiss Army Knife**
Multi-tool for diverse tasks: formatting, analysis, writing, brainstorming.

🔭 **The Visionary**
Strategic thinking and big-picture analysis. Best for planning and initiative design.

Create your own persona with custom instructions — e.g., a "DC Operations Manager" with your site's SOPs built in.

Access personas from the sidebar. They shape HOW the AI responds — same model, different behavior.

Prompt Engineering Basics

15 minutes

50% Your Prompt Is Half the Equation

MIT researchers found that half of AI performance comes from the model — and the other half from **how you write your prompt**. The best prompters weren't engineers. They were people who expressed ideas clearly.

MIT Sloan School of Management, 2024

WITHOUT PROMPT SKILLS

83% hallucination rate — AI confidently makes up facts when given vague instructions
Mount Sinai School of Medicine

78% of AI project failures caused by poor human-AI communication, not bad technology
Enterprise AI research, 2024

5.8× performance variance — the same model gives wildly different answers depending on wording
arXiv, Llama-2 study

WITH PROMPT SKILLS

21% hallucination rate — a 4× reduction just from writing clearer, structured prompts
Mount Sinai School of Medicine

40% higher quality output when users apply a structured prompting framework
Harvard / BCG, 758 consultants

34% productivity boost for beginners — new AI users gain the most from learning to prompt well
Stanford / MIT, 5,000 workers

Bottom line: Bottle Rocket already has the best AI models. The gap between a frustrating AI experience and a powerful one is **your prompt**.

The Prompt Formula: ROLE + CONTEXT + TASK + FORMAT

ROLE + **CONTEXT** + **TASK** + **FORMAT**

ROLE
Tell the AI who to be

"You are an experienced operations manager at a distribution center"

CONTEXT
Give background information

"I manage the morning shift at Tampa DC with 45 associates"

TASK
State exactly what you need

"Create a standup summary for today's meeting"

FORMAT
Specify the output structure

"Bullet points, under 150 words, 3 sections"

Good Prompt vs. Bad Prompt

Vague Prompt

Help me with standup prep

Generic, unfocused, not usable without heavy editing. Missing: role, context, specific task, format constraints.

Structured Prompt

ROLE: Experienced ops manager at a Tesla DC.

CONTEXT: Tampa DC, morning shift, 42 associates, 2,847 orders yesterday, conveyor issue 2-4 PM.

TASK: Standup summary — yesterday's performance and today's priorities.

FORMAT: Bullets, under 150 words, 3 sections.

Specific, actionable, ready to use with minor edits.

The same AI model produced both outputs. The difference is entirely in the prompt.

Your AI Dashboard — Bottle Rocket
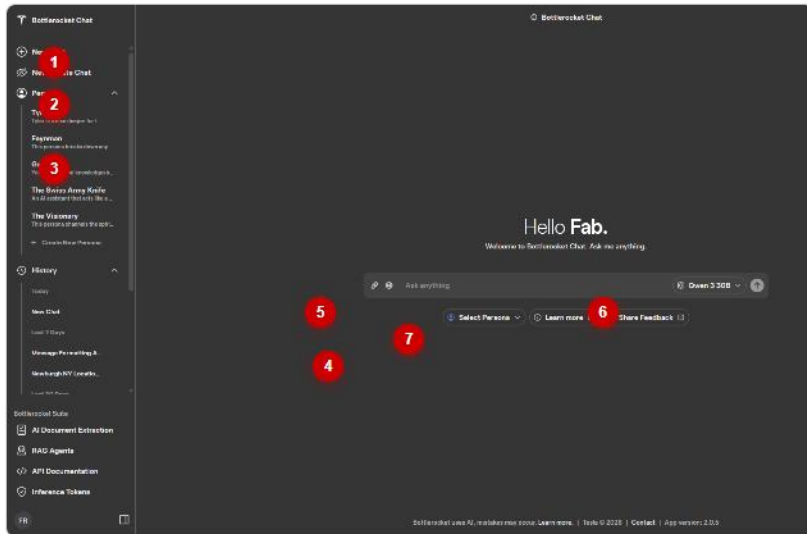
1. **New Chat** — Start fresh conversation
2. **Private Chat** — Not saved to history
3. **Personas** — Change AI personality
4. **File Upload** — Attach PDFs, data files
5. **Web Search** — Real-time info toggle
6. **Model Selector** — Choose AI model
7. **Select Persona** — Quick persona switch

All models inside BR are **approved for Tesla Data**

Explore Every Feature

Open go.tesla.com/chat and try each feature. Click on everything. You can't break it.

01
Switch Models

Click the model name (bottom-right). Try Qwen 3 30B, then switch to Grok 4. Ask the same question to both.

✓ Compare the two answers

02
Try a Persona

Click "Select Persona" below the chat. Pick Feynman. Ask it to explain something complex from your DC.

✓ Notice the tone change

03
Toggle Web Search

Click the globe icon. Ask "What's the weather at my DC today?" Compare with web search ON vs OFF.

✓ See real-time data appear

04
Upload a File

Click the paperclip icon. Upload any PDF or document. Ask BR to summarize it in 3 bullet points.

✓ Get a summary back

05
Start Private Chat

Click "New Private Chat" in the sidebar. This mode doesn't save history. Try a quick question.

✓ Check it's not in history

06
Browse Personas

Open the Personas section in the sidebar. Read each persona description. Pick your favorite for daily use.

✓ Pick your go-to persona

10 MIN Explore freely. Raise your hand if you find something cool to share with the group.

First Guided Prompt — Standup Summary

1. Open Bottle Rocket (go.tesla.com/chat)
2. Select Qwen 3 30B model (the default — fast and reliable)
3. Copy the prompt template (provided)
4. Replace placeholders with your DC's real information
5. Submit and review the output
6. Check: Is it under 150 words? In bullet format? Relevant?

20
MINUTES
5 min fill | 10 min test | 5 min share

Tips: Add specific details. If output is too generic, add more CONTEXT. If format is wrong, be more explicit in FORMAT.

EXERCISE 1 **Your Prompt Template**

ROLE: You are an experienced operations manager at a Tesla distribution center. CONTEXT: I manage the [SHIFT] shift at the [YOUR CITY] DC. Yesterday we processed approximately [NUMBER] orders with [NUMBER] associates. [Describe one thing that happened yesterday]. TASK: Create a brief standup summary highlighting yesterday's performance and today's priorities. FORMAT: Use bullet points. Keep it under 150 words. Include: 1) Yesterday's key metrics 2) Key highlight or challenge 3) Today's focus areas

TIP

Replace every [PLACEHOLDER] with your real DC information before pasting into Bottle Rocket.

Time: 20 min   **Tool:** Bottle Rocket

Model Showdown — Same Prompt, Different Brains

1 Copy the prompt below into Bottle Rocket with **Qwen 3 30B** selected

2 Read the output. Then switch to **Grok 4** and paste the exact same prompt

3 Try a third model: **Gemini 2.5 Pro**. Compare all three outputs side-by-side

4 **Discuss with your table:** Which was fastest? Most detailed? Most useful for YOUR work?

USE THIS PROMPT

"You are an operations analyst at a Tesla distribution center. Analyze the following scenario: our daily shipment volume increased 40% this week but headcount stayed the same. Recommend 3 specific actions I can take today to maintain delivery targets. Format as a numbered list with estimated impact."
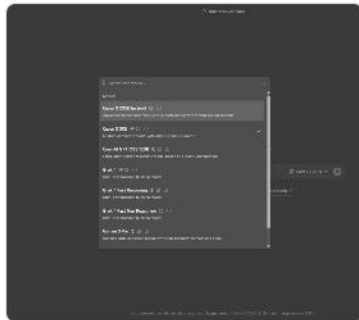
15
MINUTES

Notice: Speed vs. depth trade-off. Qwen 3 30B is fast. Grok 4 thinks deeper but takes longer.

Operations Prompts + Persona Comparison

1. Choose 3 real work tasks (email, meeting prep, performance summary)
2. Write each using ROLE + CONTEXT + TASK + FORMAT
3. Test each in Bottle Rocket and review results
4. For ONE prompt: switch to the Feynman persona and re-run
5. Compare persona output vs. no-persona output

25
MINUTES
5 min topics | 15 min test | 5 min persona

Tips: Use tasks you actually do every day. Notice how the persona changes tone and structure.

CARD A
Email Drafts
- Vendor late delivery
- Leadership update
- Team recognition
- Customer escalation

CARD B
Meeting Prep
- Safety meeting agenda
- 1:1 with supervisor
- Shift handoff debrief
- QBR talking points

CARD C
Performance Summaries
- Weekly metrics
- Incident report
- Shift handoff
- Regional highlight email

CARD D
Analysis & Problem-Solving
- Missed shipments
- Overtime spikes
- Dock door efficiency
- Retention brainstorm

Pick 3 tasks. Write an RCTF prompt for each. Test in Bottle Rocket. For **ONE** prompt, switch to Feynman persona and compare.

25 MIN

Three Tiers of AI: Chatbot → Workflow → Agent

Not all AI is the same. Understanding the tiers helps you pick the right tool for the job.

1 AI Chatbot

Answers questions by retrieving content

Fetches information from a **known set** of documents or data. Matches your question to stored content and generates a response.

→ Retrieves from documents (RAG)
→ Matches query to content
→ Generates formatted answer
→ **Read-only** — cannot take action

TESLA EXAMPLE
Bottle Rocket Chat — ask questions, get answers from AI models. **Nabu agents** with uploaded docs.

2 AI Workflow

Answers questions by invoking systems

Orchestrates data from **multiple sources** or services. Goes beyond static retrieval — calls APIs and combines results.

→ Calls external APIs for live data
→ Combines multiple data sources
→ Computes or aggregates results
→ **Read-only** — pulls data, doesn't change it

TESLA EXAMPLE
IT Assist / Employee Assist — pulls from HR systems, knowledge bases, and ticketing to answer your question.

3 AI Agent

Performs actions by invoking systems

Autonomous entity that **executes tasks** and influences the environment. Makes decisions and acts on them.

→ Perceives situation & decides actions
→ Calls APIs to read AND write
→ Automates workflows end-to-end
→ Monitors outcomes, adjusts behavior

TESLA EXAMPLE
Future state — an agent that monitors Order Monitor, auto-escalates stuck VOR orders, and reassigns pickers.

Tesla Employees Are Already Building AI Agents

5-Why Analysis Agent
OPERATIONS

Guides structured root cause analysis for production issues

Parts SOP Generator
OPERATIONS

Generates standard operating procedures from parts documentation

Watt Wizard
ENERGY

Answers energy storage and battery technical questions

SPARQ
ENGINEERING

Spark plug analysis and quality inspection checks

Megafactory 5S
MANUFACTURING

5S audit assistance for factory floor operations

Your Agent?
YOUR DC

What repetitive task could AI handle for your team?

Built on Nabu (nabu.bottlerocket.tesla.com) — 38 agents and growing. No code required.

Nabu — Build Your Own AI Agent

1 **Create Agent** — Start building in 1 click

2 **Learn More** — Docs and tutorials

3 **Search Agents** — Find existing agents

4 **Agent List** — Live agents with status badges

5 **Descriptions** — What each agent does

38+

agents built by Tesla employees
No code required. Just markdown + prompts.

Open: nabu.bottlerocket.tesla.com — v2.1.1

The Prompt Progression — Same Data, 4 Prompts

Same AI model. Same data. Only the prompt quality changes.

1. Trainer opens 4 Nabu agents (Order Monitor Levels 1–4), all loaded with same 1,000-order dataset

2. Same question asked to all 4: "What are the most urgent orders I should worry about?"

3. Watch outputs side-by-side — Level 1 (generic) → Level 4 (expert with part numbers and actions)

4. Group discussion: What made the difference at each level?

25
MINUTES
5 min setup | 10 min demo | 10 min discuss

Key insight: The AI model is the same. The data is the same. The ONLY variable is the quality of your instructions.

The 4 Prompt Levels ACTIVITY C

Level 1 Analyze the file.

The Lazy Prompt
3 words, zero guidance

Level 2 You are a warehouse analyst. Look at the order data and tell me what's going on. Find any problems and give me a summary. Use bullet points.

Some Effort

Level 3 ROLE: Operations analyst… CONTEXT: Order Monitor export, Hot Calls, VOR…
TASK: Executive summary, flag urgent, identify problems, 3-5 recommendations…
FORMAT: Headers, bullets, bold numbers

Good Structure
(RCTF Framework)

Level 4 Everything from Level 3 PLUS: 50+ field definitions, status lifecycle, pick location format, part groups, route types, analysis thresholds

Expert

Same AI model. Same data. Only the prompt changes.

Design Your Agent on Nabu

1. Open **nabu.bottlerocket.tesla.com**

2. Browse published agents — note 5-Why Analysis and Watt Wizard as examples

3. Click **"Create New Agent"**

4. Name your agent: **[Your DC] + [Use Case]** (e.g., "Tampa Shift Handoff Agent")

5. Write a system prompt using **ROLE + CONTEXT + TASK + FORMAT**

6. Upload one document to the knowledge base (an SOP, checklist, or template)

7. Test in the **Playground** with a sample query

20
MINUTES
5 min explore | 10 min build | 5 min test

AGENT IDEAS

Shift Handoff Agent — Generates end-of-shift summaries

New Hire Onboarding — Answers associate FAQs

Safety Incident Logger — Structures incident reports

Start simple: One task, one document. You can improve it in Week 3.

Bottle Rocket
Deep Workshop

30 minutes

Bottle Rocket Deep Workshop

A
Analyze Headcount

Analyze headcount allocation across shifts

FILE UPLOAD

B
Research Topic

Research a real-world operations topic with live data

WEB SEARCH

C
Vendor Communication

Draft vendor communication for delayed shipment

MODEL COMPARISON

D
Training Outline

Create training outline for new associates

FEYNMAN PERSONA

---

30 MIN   Complete at least 2 scenarios. Stretch goal: all 4.

Exercise 4: Scenario Data `EXERCISE 4`

Scenario A: Headcount `FILE UPLOAD`

```
SHIFT DATA - WEEK OF JAN 20-24:
Mon:
Morning 18 assoc / 1,247 orders
Afternoon 22 / 1,893
Night 8 / 412

Tue:
Morning 17 / 1,189
Afternoon 23 / 2,104
Night 9 / 387

Wed:
Morning 19 / 1,302
Afternoon 21 / 1,756
Night 7 / 445

Thu:
Morning 18 / 1,278
Afternoon 24 / 2,231
Night 8 / 398

Fri:
Morning 20 / 1,534
Afternoon 22 / 2,087
Night 6 / 356
```

Scenario B: Safety Notes `WEB SEARCH`

```
RAW INCIDENT NOTES:
- happened around 10:15am near
  dock door 7

- forklift operator Carlos M.,
  badge #4472

- backing up with loaded pallet,
  didn't see pedestrian walkway
  occupied

- near-miss, nobody hurt

- pedestrian was new hire Jamie L.
  - no high-vis vest

- vests on backorder, last 3
  classes didn't get them

- backup alarm on forklift #12
  "kind of quiet" for a week

- told maintenance to check
  alarm TODAY
```

Hallucination Detection & Verification

Hallucination: AI confidently presents made-up facts as truth

Always cross-reference specific claims against known data

Ask the AI: "What sources support this?" — then verify those sources exist

The 80/20 trust rule: trust formatting and structure, verify facts and numbers

Never forward AI-generated numbers to leadership without checking them

You own the accuracy of everything you send — AI assisted, you are accountable

Trust the draft. Verify the data.

Hallucination Hunt — Spot the Errors

1. Read the AI-generated shift handoff summary (projected on screen)
2. The output looks professional and confident — but contains 3 deliberate errors
3. Find all 3 errors: a made-up statistic, an incorrect OSHA claim, and fabricated internal data
4. Write down what's wrong and why it matters
5. Compare with answer key — discuss as a group

10
MINUTES
5 min hunt | 5 min debrief

Warning: AI hallucinations are dangerous BECAUSE they sound confident. Always verify statistics, regulatory claims, and internal data.

Hallucination Hunt: Find the 3 Errors **EXERCISE 5**

● AI-GENERATED OUTPUT

Shift Handoff Best Practices for Tesla Distribution Centers

Effective shift handoffs are critical to maintaining throughput and safety. According to the **Tesla Operations Handbook (Section 12.4)**, all shift handoffs must include documented transfer of open issues, pending orders, and equipment status.

Key Metrics: Orders processed vs. target (95% fulfillment rate), current backlog, equipment status, staffing notes.

Safety Protocol: Tesla's Global Safety Standard requires outgoing shift leads complete a safety walkthrough. This includes verifying all forklift operators have completed their **OSHA-mandated daily 15-minute pre-shift vehicle inspection checklist**.

Communication Standard: According to a **2025 MIT Supply Chain study**, DCs with structured handoff protocols saw a **47% reduction in errors**. Tesla's own internal data shows DCs using the standard handoff template **reduced missed shipments by 31% in Q3 2025**.

Error 1:
_ _ _ _ _ _

Error 2:
_ _ _ _ _ _

Error 3:
_ _ _ _ _ _

Hallucination Hunt: Answer Key ANSWER KEY

ERROR 1 **Made-Up Statistic**

The "2025 MIT Supply Chain study" with a "47% reduction" does not exist. AI fabricated the source, year, and number.

HOW TO CATCH: *Search for the cited source*

ERROR 2 **Incorrect Regulation**

OSHA does not mandate a "daily 15-minute pre-shift inspection." Standard 1910.178(q)(7) requires pre-use examination but specifies no time requirement.

HOW TO CATCH: *Verify against actual regulation*

ERROR 3 **Fabricated Internal Data**

AI cannot access Tesla internal data. The "31% reduction in Q3 2025" is invented. Bottle Rocket has no access to operational databases.

HOW TO CATCH: *Did I provide this data, or did AI invent it?*

If you didn't provide the data and can't verify the claim — don't use it.

Iterate, Verify & Improve

INSTRUCTIONS

1. Take your best prompt from Exercise 2

2. Review the output critically — what is not quite right?

3. Apply a verification technique (cross-reference a claim, check a number)

4. Use a follow-up prompt to improve (e.g., "Make this more concise")

5. Bonus: Re-run with web search ON and compare accuracy

25
MINUTES
5 min review | 10 min verify & iterate | 10 min share

Tips: Don't rewrite the entire prompt — just tell the AI what to fix. Iteration is the skill that separates good AI users from great ones.

Iteration Techniques

WHAT'S WRONG	FOLLOW-UP PROMPT TO TYPE
Too long	*"Shorten this to under 100 words. Keep the 3 most important points."*
Too short	*"Expand the section on [topic]. Add specific examples."*
Too formal	*"Rewrite in a more direct, conversational tone."*
Too casual	*"Make this more professional. This is going to the regional director."*
Missing section	*"Add a section at the end with 3 specific next steps and owners."*
Wrong format	*"Convert from paragraphs to bullet-point summary. Group by category."*
Too generic	*"Make this specific to a Tesla DC processing 3,000+ orders/day."*
Unverifiable claims	*"Remove statistics you can't verify. Replace with [INSERT ACTUAL METRIC]."*

PRO TIP
Don't rewrite the entire prompt — just tell the AI what to fix. Iteration is the skill that separates good AI users from great ones.

Time Savings Baseline — Measure Your Impact

1. Choose a tracking method: spreadsheet, notebook, or sticky note
2. Create the tracking template: Date | Task | Manual Time | AI Time | Saved | Notes
3. Measure your first baseline from today's exercises (pick one task)
4. Set your Week 2 goal: "I will use BR for ___ tasks/day, targeting ___ min saved"
5. Pick a daily tracking trigger (e.g., "every time I close BR, I log it")

15
MINUTES
3 min setup | 7 min measure | 5 min goals

Targets: Week 1: 15–20 min/day saved. Week 4: 40–60 min/day. Over a year at 40 min/day = 173 hours saved.

EXERCISE 7 **Time Savings Tracker**

DATE	TASK DESCRIPTION	MANUAL (MIN)	AI (MIN)	SAVED (MIN)
Jan 31	Standup talking points	15	3	12
Jan 31	Vendor email (late delivery)	12	4	8
Jan 31	Safety incident summary	20	5	15

SET YOUR GOAL

Cautious	Confident	All-In
1–2 tasks/day	2–3 tasks/day	3–5 tasks/day
10–15 min saved	15–25 min saved	20–30 min saved

Your Prompt Library — Quick Reference

1	**Morning Standup** Qwen 3 30B	15 min
2	**Shift Handoff** Qwen 3 30B	12 min
3	**Overnight Briefing** Grok 4 Mini	10 min
4	**Vendor Escalation** Qwen 3 30B	12 min
5	**Leadership Update** Qwen 3 30B	15 min
6	**Team Announcement** Grok 4 Mini	8 min

7	**Weekly KPI Summary** Grok 4	20 min
8	**OT Root Cause** Grok 4	25 min
9	**Safety Meeting Agenda** Qwen 3 30B	10 min
10	**1:1 Supervisor Prep** Grok 4 Mini	8 min
11	**Safety Incident Report** Qwen 3 30B	20 min
12	**Onboarding Checklist** Qwen 3 30B	15 min

Numbers/analysis → **Grok 4**   Writing/reports → **Qwen 3 30B**   Quick tasks → **Grok 4 Mini**

Week 1 Key Takeaways

01 AI augments your expertise — it handles repetitive work while you focus on judgment and strategy

02 Policy is non-negotiable — Bottle Rocket for Tesla Data, the Two-Question Test is your guide, platform determines protection

03 ROLE + CONTEXT + TASK + FORMAT — the four-part formula transforms vague requests into actionable outputs

04 Iterate, verify, and use personas — your first output is a draft, not a deliverable. Check facts, refine, then use.